



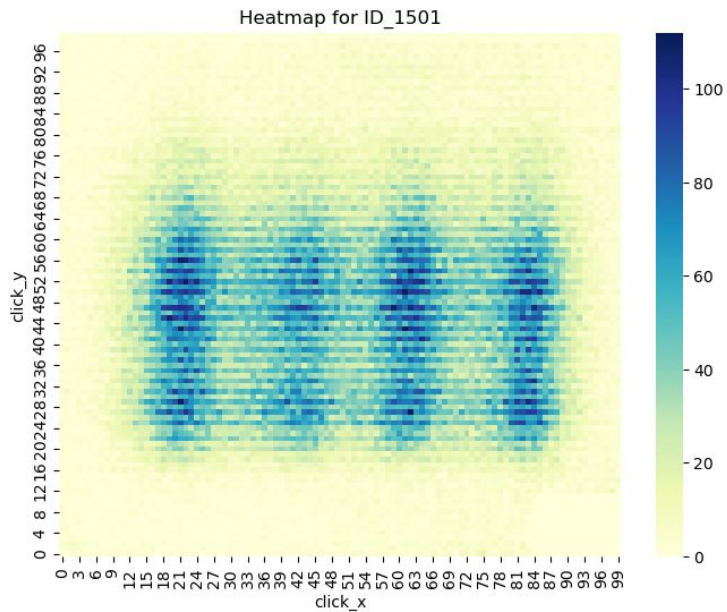
Capstone Project: Heatmap Anomaly Detection

Week 10 Progress Report

This week:



1. (from last week) Fine-tuned ResNet on different tasks
 - a. Applied to new data
2. Test models on additional dataset

Questions



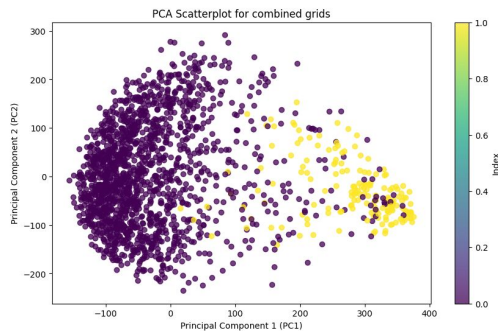

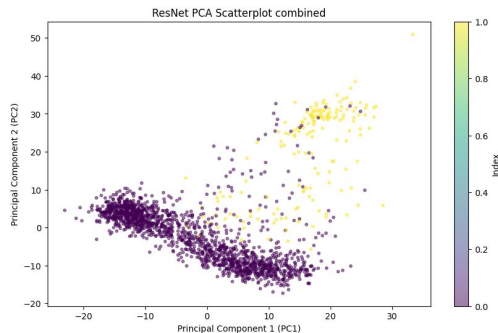

- New dataset:
 - 2 grid_ids: One seems to only have 4 products?
 - Confirmation: $\text{click_x_rel}/\text{click_y_rel} = 40 \times 40$ binned?
 - Any particular reason?
- Clarifications about scale:
 - “25 M banners/day”:
 - How many new clicks/banner/day?
 - How is the data formatted?
- Compute power seems quite low (esp 8GB memory)?
 - How much time could we use / day?
 - Problematic even for inference.
- Confirmation:
 - We have 5 grid_id's
 - The triple (grid_id, #products, domain) =: banner?

Pretrained ViT/ResNet (recap):




- Feed (transformed and binned) heatmaps into pre-trained ViT/ResNet.
 - google/ViT: Transformer-based architecture, 14M images (224x224), 21k classes
 - Microsoft/ResNet-1k: trained on ImageNet-1k (224x224), 1k classes.
- Extract features (before classification head)
 - ViT → 151296 dim'l feature vector
 - ResNet → 2048 dim'l feature vector
- Play with upsampling (bootstrapping + noise)
- PCA and other dim'l reduction techniques
 - Apply clustering methods


Data enhancement (recap):



- We now bootstrap upsample to 50'000 clicks per image.
- PCA's look a lot more promising.
- MVU also look very indicative of deep structure from pre-trained models.
- => continue with enhanced data from now on.

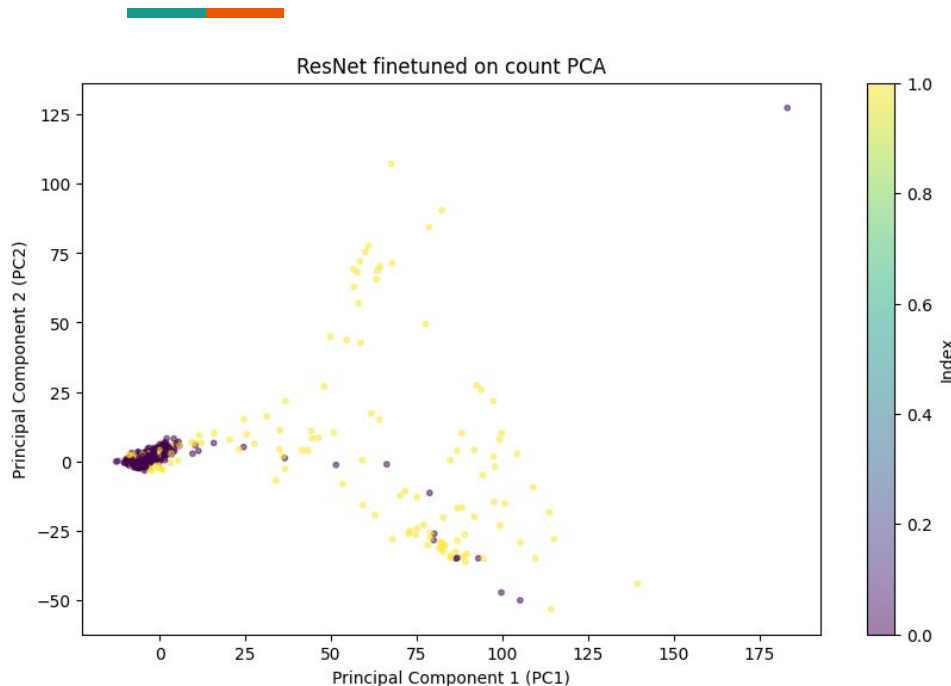


Cosine similarity (recap)




- PCA → gives some intuition, but very limited
- MVU → shows there is underlying structure.
- Feature vectors usually evaluated using **cosine similarity** (angle between vectors in high-dimensional feature space)
- **Cosine similarity** has actual meaning (while PCA does not really, since “maximum variance” is inherently imprecise).
- Draw picture of cosine similarity using ResNet and ViT feature vectors.
- What choice of “anker”?
 - Pick random choice
 - Pick average non-broken grid_id's.

Fine-tune ResNet-50 on number of clusters




- Fine-tune feature vectors of ResNet-50 architecture:
 - Pretrained ResNet
 - Generate synthetic images:
 - Random # of clusters (0-20)
 - Random center for clusters
 - Random covariance for Gaussians
 - Random # of clicks/cluster
 - Overlay random noise
 - Add classification head (softmax)
 - “Learn” detection of number of clusters (cross-entropy loss).
- After some training get ~ 40% accuracy on predicting # of clusters from a given image (somewhat depending on some of the hyperparameters).
- PCA combines the two clusters into one (= have the same number of clusters).

Fine-tune ResNet-50 on position of clusters



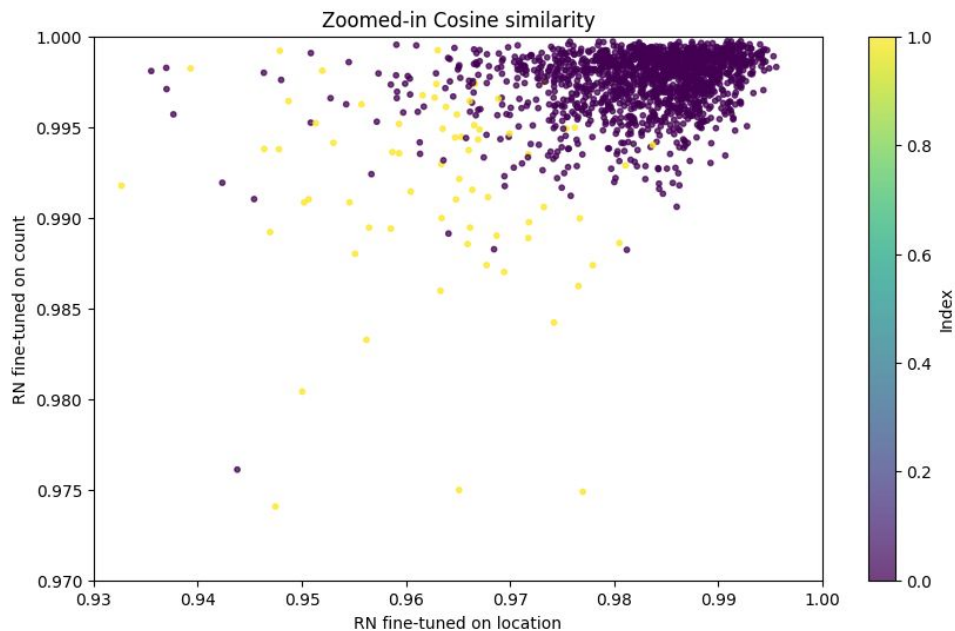
- Fine-tune feature vectors of ResNet-50 architecture:
 - Pretrained ResNet
 - Generate synthetic images:
 - Random # of clusters (0-6)
 - Random center for clusters
 - Random covariance for Gaussians
 - Random # of clicks/cluster
 - Overlay random noise
 - Add classification head providing a set of centers $\{(x_1, y_1), (x_2, y_2), \dots, (x_6, y_6)\}$
 - Setwise distance loss function; **Chamfer distance**:
$$c(A, B) = \sum_{i=1}^{|A|} \text{dist}(a_i, \text{nn}(B, a_i)) + \sum_{j=1}^{|B|} \text{dist}(b_j, \text{nn}(A, b_j))$$
- PCA shows two clusters “pointing” in different directions (=captures difference in location).

Cosine Similarity Count vs location





- Anker = mean of 2 unbroken clusters from each grid_id
- Cosine similarity shows a clear distinction between broken and non-broken banners
- Blue dots outside of central cluster mostly misclassified.

Cosine Similarity Count vs location (zoomed in)





- Anker = mean of 2 unbroken clusters from each grid_id
- Cosine similarity shows a clear distinction between broken and non-broken banners
- Blue dots outside of central cluster mostly misclassified.


Cosine Similarity Count/loc vs Pretrained



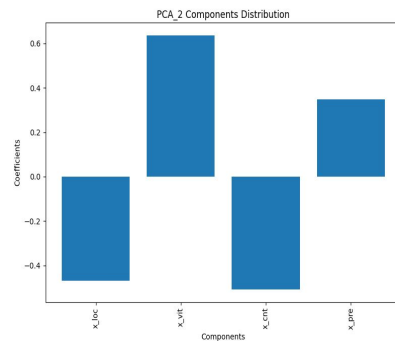

Cosine Similarity Count/loc vs ViT




ViT, ResNet, fine-tuned Resnet



- Have now 4 powerful features (pre-trained ViT/ResNet, fine-tuned ResNet on cluster location/count).
- How do we best combine this.

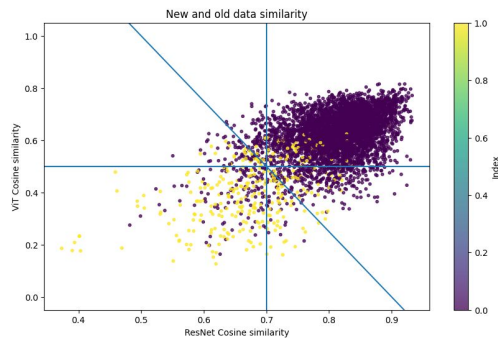
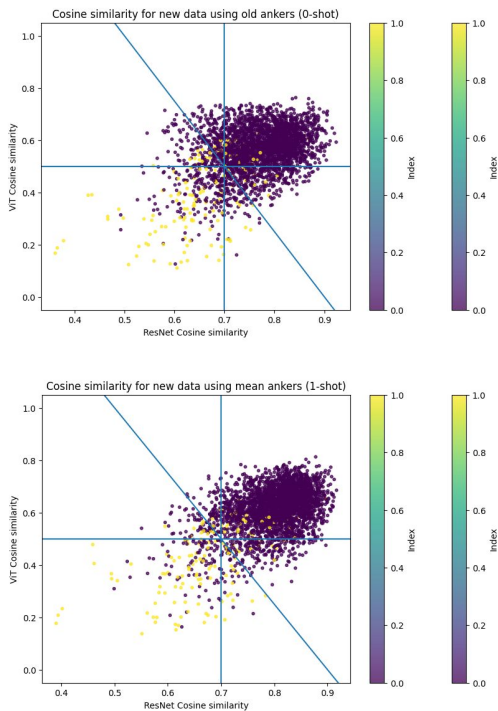


Pretrained ViT/ResNet for new data:




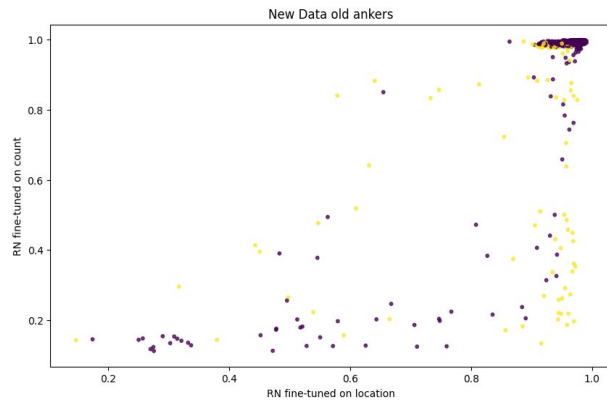
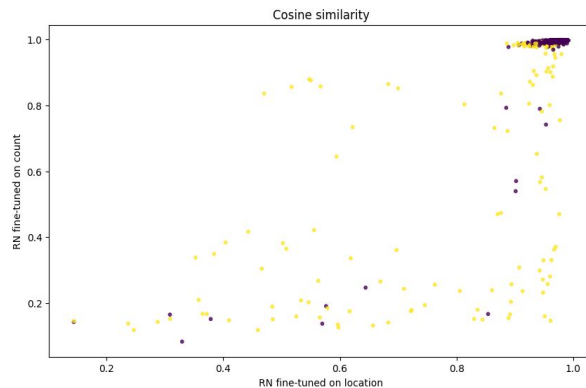
- Apply these methods for new dataset.
- Notice that we've identified 1 (at least) image as broken in previous data that looked like first image.
- Banner 333346 (top) harder to see structure → maybe change noisy upsampling?

Pretrained ViT/ResNet for new data:






- 0-shot and 1-shot performance for new data.
 - “Old” threshold still very good
 - 1-shot better than 0-shot
- Combined data looks quite good too.
- Need to analyse some of the “misclassifications”. (to do)



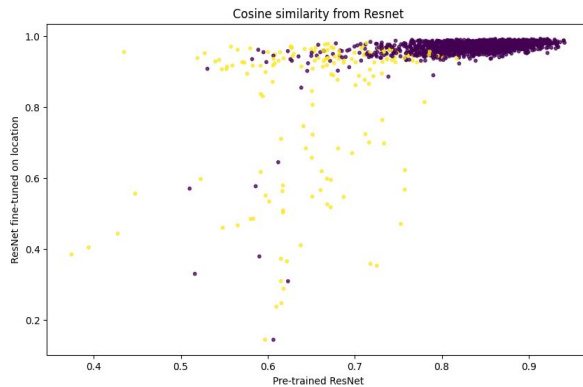
O/1-shot: count vs location






O/1-shot: count vs pretrained






O/1-shot: location vs pretrained



0/1-shot: ViT vs location



O/1-shot: ViT vs count




Next steps:




- Further investigate 0/1-shot performance on new dataset.
- Converge towards actual predictor
- Combine 4 ResNet/ViT features into a meaningful predictor.
 - Are there other synthetic tasks?
 - Combine with metrics features and PCA?
- SimCLR -> get it running
- Measure compute requirements more carefully:
 - Data pre-processing slow (I think can be sped up)
 - Inference times slow-ish → can be parallelized (xCPU-number speedup)

New data









- Exploration
- Implementation



Clearly broken banners = $74/2061 = 3.59\%$



Clearly broken banners = $86/969 = 8.88\%$





Test on subSet of New Data

- Use our trained models from previous data
- Predict on subset of the new data (796)
- 550(333519) & 246(333346)

Predicted Result from previous combined model:

- Anomaly Percentage: 7.79%
- Number of Broken Banner Predicted: 62

On new data

SVM:

F1 Score: 0.14

KNN:

F1 Score: 0.39

K-Means:

F1 score: 0.02

DBScan:

F1 Score: 0.49

Isolation Forest:

F1 Score: 0.37

Combined Model:

F1 Score: 0.40

On previous data(test set)

SVM:

F1 Score: 0.25

KNN:

F1 Score: 0.83

K-Means:

F1 score: 0.24

DBScan:

F1 Score: 0.72

Isolation Forest:

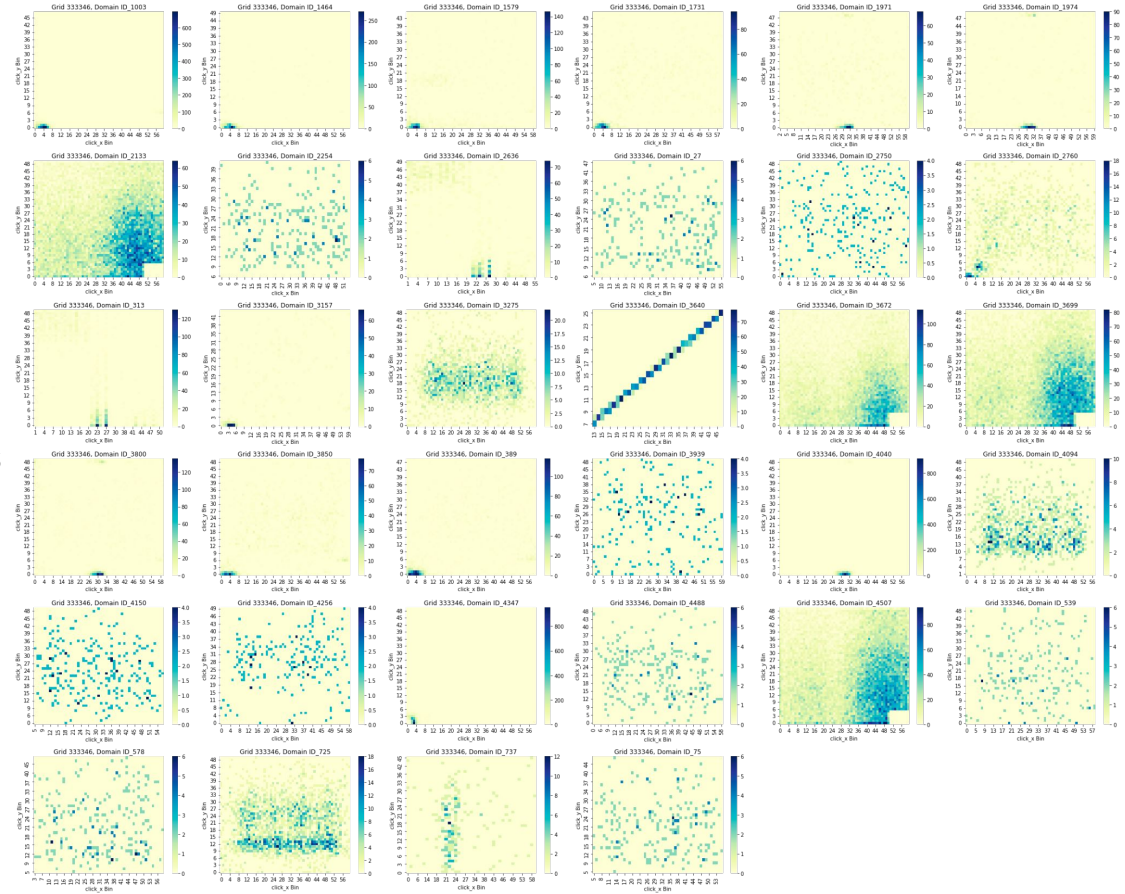
F1 Score: 0.85

Combined Model:

F1 Score: 0.87

Combined Model Result1:

- This is Grid ID 333346
- For Grid ID 333346, there are 34 anomalies
- $34/246 = 0.138 \rightarrow 13.8\%$

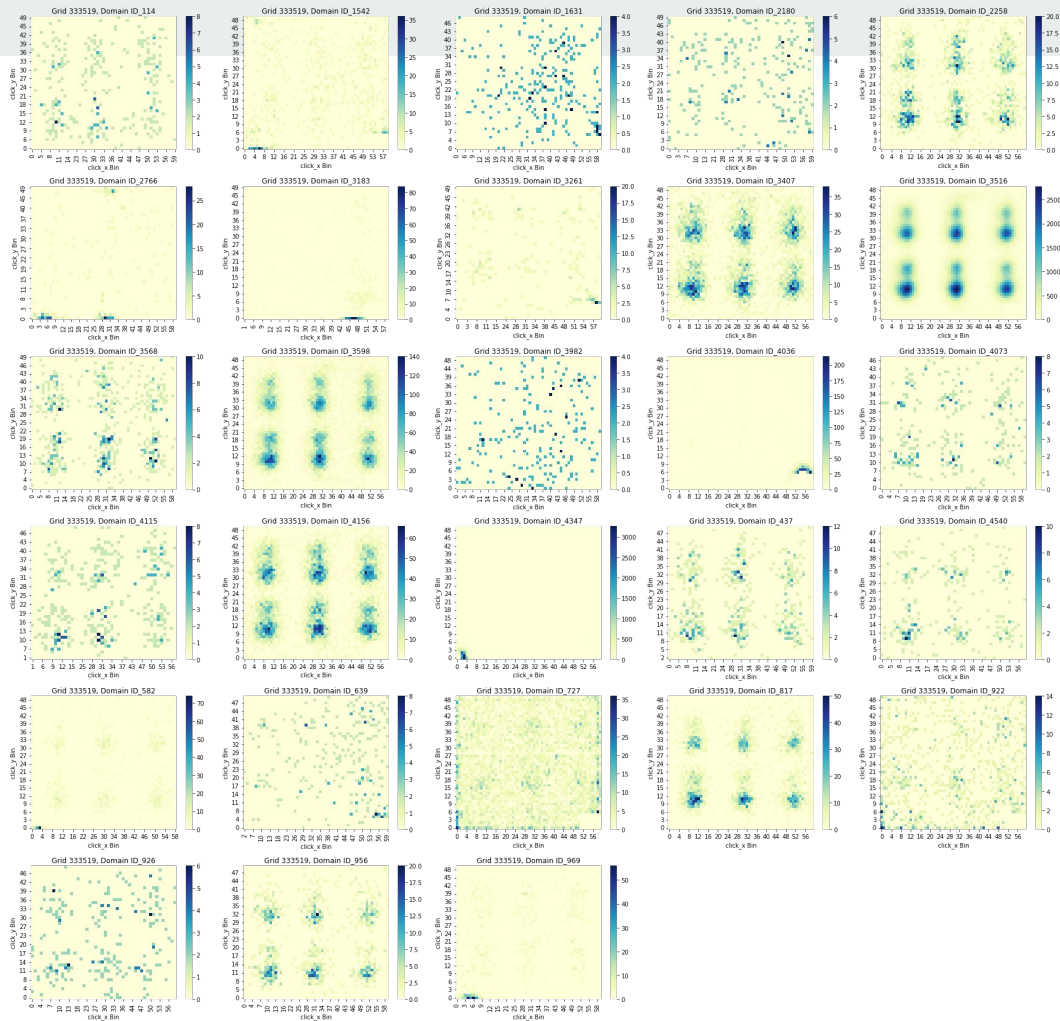


Combined Model Result2:


- This is Grid ID 333519
- For Grid ID 333519, there are 28 anomalies
- $28/550 = 0.051 \rightarrow 5.1\%$

Further steps:

- More exploration about the data



Old/New Statistical approach (wip):



- Cumulative clicks → empirical distribution f
- New banner → drawn from distribution
- Draw probability:
 - Can compute $P[\text{New banner} | \text{emp distribution}]$
 - Normalize: $p = P[\text{NB} | \text{ED}] / P[\text{Avg B} | \text{ED}]$
 - If $p < \text{threshold}$: broken.
 - ~30% identified as broken
- Chi-squared:
 - Chi-squared test:
 - Is the underlying distribution the same between ED and NB?
 - ~30% identified as broken
- LRT:
 - Estimate “bad banner” distribution
 - Compute $p = P[\text{NB} | \text{ED}] / P[\text{Avg B} | \text{BB}]$
 - Likelihood-ratio-test (to do).
- Upshot:
 - Lightweight and can be easily implemented on data-stream.
- To do:
 - Different resolutions
 - Upsample/downsample
 - “pseudo-distance” between good and bad banners (e.g. KL-divergence)



Appendix